

A Feature-Generalizable Technique for Neural Conditioning

Jack D.V. Carson

Under the Direction of
Dr. Alan Jasanoff, Kevin Chung, and Dr. Ying Jiang

Research Science Institute
July 28, 2023

Abstract

Advances in brain-computer interface technology seek to enable nuanced control of neural activity patterns. This study demonstrates a framework combining artificial intelligence and high-resolution optical sensing to identify and modulate intricate signaling dynamics. Genetically-encoded calcium indicators provide fluorescent readouts of neuronal firing. An autoencoder neural network compresses these optical data into a compact latent space, extracting interpretable features. The custom software Rasa coordinates data flow, deploys the models to translate neural sequences into vector representations. By comparing embedding vectors of target and observed activity, Rasa identifies desired patterns and administers neurofeedback accordingly. Initial *in vivo* validation demonstrates increased activity of simple correlated firing patterns under this novel stimulation paradigm versus unmodulated recordings in rat motor cortices. While limited by experimental constraints, these preliminary results highlight the potential of integrated machine learning techniques and fine-grained optical sensing to reinforce complex behavior. Looking forward, such AI systems could unlock new therapeutic abilities to remedy dysfunctional neurological signaling underlying disease states.

Summary

This research investigates brain-computer interaction, making it possible to train complex patterns of brain activity much more easily. Training begins with GCaMP protein-based sensors that read high resolution brain signals on a microscopic region of neural activity. This paper presents a new way of training to specific brain signals by comparing the weights of a machine learning model, which could allow us to control more complex signals and speed up training. Furthermore, we developed a new software package titled Rasa to address the issues of data delay, reliability, and performance that come up with the high throughput of neural sequence data. When we bring Rasa together with our new training technique, it offers a well-rounded solution to the limitations we currently face in brain-computer interaction.

1 Introduction

As scientific understanding of functional neural activity increases, researchers have begun to take an interest in the treatment of the human brain as an input-output (I/O) computational machine [1]. This interest arises from the neurological capacity for conditioning that the brain uses in adapting to its physical environment. A major question in the field of neuroengineering is to what extent this plasticity can be controlled by electrical signals from an external source, which could effect conditioning through the already-observed mechanisms; the impact of the external signals can be determined by a real-time analysis of the brain's output data streams.

Current neuroimaging and neurophysiological sensors, such as electroencephalography (EEG) and electrocorticography (ECoG), expand our ability to model brain activity, but both technologies fall short in capturing delicate electrical signals with spatial precision due to factors such as large electrode size, which obscures spatial information of signals, and non-invasive application, which limits the precision and sensitivity of output streams [2] [3]. Furthermore, both ECoG and EEG confound hemodynamic methods of analysis such as fMRI, a valuable tool for mapping white and grey matter shifts in the brain that indicate a response to training. To address this, calcium sensors based on genetically embedded fluorescent proteins such as GCaMP have emerged as a potential alternative for capturing neurological activity at finer scales [4]. The data generated from calcium sensors can form the foundation of a neurofeedback output stream.

In the wake of recent advancements in neuroscience, there is also a growing consensus that electrical deep brain stimulation, apart from acting as a data receiver from neural sensors, can potentially serve as a robust input channel for neurofeedback and thus form an integral component of a Brain-Machine Interface (BMI). This process commences with an observation of neural activity, which is subjected to online analysis to provide a requisite

stimulation [5] [6] [7]. However, the current neurofeedback techniques are bound by both the complexity of the conditionable signals and also the experimental difficulty of conditioning, which together conspire to severely inhibit the encoding of explicit brain signals and confine conditioning to simplistic binary-enforceable signals (i.e., Do we observe a spike?: Yes/No).

This study seeks to challenge these limitations, targeting the complexity of encodable signals and the rate at which these signals can be trained. We propose to optimize the training of neural signals by defining a continuous function that can be used to robustly evaluate the closeness of two arbitrary signals. We use an autoencoder neural model to define a vector representation of a particular neural signal as the values of latent variables at the bottleneck of an autoencoder model. Through comparing the distance of a target and active vector in the latent space of an autoencoder model, we were able to define an optimizable point to approach for conditioning. This technique aims to amplify the diversity and complexity of encodable signals within the existing feedback framework, while also enhancing the precision and speed of training for even simple signals. In tandem with this, we introduce a new software package, developed in the Rust programming language and specifically crafted to address the glaring issues of data latency, integrity, and performance encountered with existing data logging and real-time visualization software. By integrating this new software with our innovative optimization technique, we aim to introduce a refined, comprehensive solution to the existing limitations in neurofeedback and conditionable signal complexity.

2 Methods

2.1 Surgical Preparation

Several adult female rats were kept in single-housed cages and maintained in accordance with guidelines set by the Institutional Animal Care and Use Committee of the Massachusetts Institute of Technology. We performed an initial surgery to implant a 200 μm -diameter optic




Figure 1: Top: Bipolar deep brain stimulation electrode: (+3.4mm AP, +1.7mm ML, 8.5mm DV); Bottom: Virus (AAV9 containing DNA encoding GCaMP8f cargo) and optic fiber: (-2mm AP, +/-3.5mm ML, 2.5mm DV)

fiber canula in the rat’s motor cortex, allowing for the reading of GCaMP8f-based photometry data. The animals were provided with isoflurane analgesic pain relief during and after the operation in addition to general anaesthetization[8]. During surgery, the genetically encoded calcium indicator GCaMP8f was introduced virally at stereotaxic coordinates of 2.0mm anterior of bregma, 3.5mm lateral, and 2.5mm ventral from bregma, which are part of a broad area of the motor cortex hypothesized to be under the control of volitional movement [9]. The selected ensembles are believed to carry information about the rat’s intentions to move; the clarity of their purpose and lack of confounding noise makes them an ideal target for feedback. Following the recovery period from this surgery, the rats were again anaesthetized for a craniotomy to insert a deep-brain electrode that stimulates the Medial Forebrain Bundle (MFB). The MFB is a dopaminergic pathway that plays a significant role in the brain’s reward system, correlating neural signals with pleasure as part of the closed-loop feedback used in conditioning such signals [10].

2.2 Dataset

A set of approximately 18 million datapoints was generated from an Arduino setup that serializes the output of a highly sensitive photosensor reading the fluorescence intensity of GCaMP proteins passed through an amplification tube. Data from the left and right hemisphere sensors were read concurrently and saved independently. The temporal component of the sensing data was preserved and stored sequentially such that the datastream could be interpreted as the continuous line of calcium concentration, itself a proxy for activity in an neuron ensemble. The full dataset represents over 30 hours of sensing under both dormant and scanning environments. Data was captured for five rats during fMRI sessions, where neural activity was stimulated indirectly by the noise of the EPI scan. The data stream from each GCaMP sensor was subdivided into discrete subsets of 2,560 values by a sliding window with no self-overlap. The size of the window was calculated to be approximately equal to the number of data points processed per second, making each subset approximately 1 second in length. Points from this data set were averaged together in groups of 40 to form subsets with 64 values per channel. The average procedure serves purposes of (a) reducing noise among the highly sensitive photometer, (b) shrinking the data such that it can be analyzed on a millisecond timescale, and (c) increasing the ease of comparing similar features with slight offsets without the use of excessive pooling layers.

Because the vertical separation between channels, as well as the gain of the signal, are parameters to the photosensor that were non-uniform over the collection of the dataset (and separated from the actual observation), normalization of the channels was essential. For some sequence \mathbf{s} as a part of stream m ,

$$\mathbf{s}_{\text{norm}} = \frac{\mathbf{s} - \text{avg}(\mathbf{s})}{\mu_m}.$$

For the purposes of analysis without introducing additional dimensionality to the input, the two averaged channels of size 64 were concatenated into a single sequence of size 128 and separated by a consistent value such that the channels maintained their relativity but were nonetheless uniformly offset, offering an efficient and simple, albeit inelegant, way to transmit data to the encoder.

2.3 Model Design

The objective of a model for the task of comparing neural signals does not necessarily belong to the familiar tasks of classification, identification, or segmentation, for which there are extensive precedents and techniques for training. Instead, the computational approach attempts to form the modelling aspect of the abstract task of trying to modulate one signal to another signal as produced by an unpredictable *in vivo* data source. A model for this task requires a method of classifying with a continuous variable exactly how close a signal is to a target signal. In theory, this would allow incremental improvements in quality, e.g., if one wished to reinforce a strong double peak. Instead of always rewarding that activity when it happens correctly, which happens to be considerably rare and unpredictable, reward more frequent, but incorrect activity, and gradually increase the threshold for correctness based on the continuous numerical evaluation for similarity. This evaluation may be derived with a neural model table to perform analysis that is (a) efficient enough to give the impression of continuity in the evaluation stream and, (b) able to define a "loss" mechanism to show how the current target should descend towards the optimization signal.

We hypothesized that a recurrent autoencoder would be an effective solution to this task. Autoencoder models take a higher-dimensional input stream such as an image or sequence, and instead of performing inference on it, shrinks it to a narrow bottleneck layer, in which each weight represents some underlying property of the input data structure [11]. The model trains the weights by attempting to reconstruct this bottleneck into the original input data

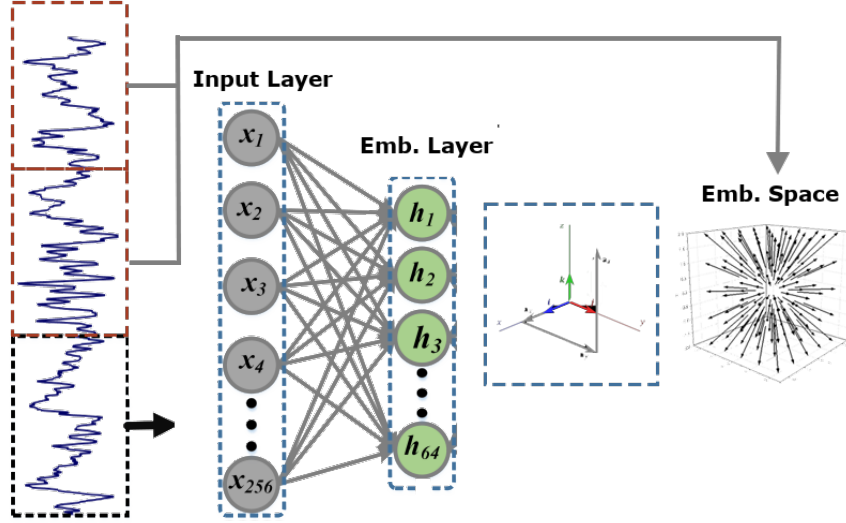


Figure 2: Visualization of headless autoencoder. Raw embedding vectors considered in embeddings space.

with no external aid. In this architecture, the interior values at the bottleneck are referred to as the *latent values* of the input. Interpreted as a vector, they form the *latent space* of the autoencoder, a inferential embedding space for the model. Most AE models use this mechanism for the purposes of generative modelling, using elaborate decoder blocks, or data compression and dimensionality reduction, querying the absolute latent values as the endpoint of the model.

The described model takes the subsets of neural activity measured through GCaMP and passes them through an autoencoder model, which represents the data in a 16-dimensional vector of latent values. Following that, an ideal target signal has been subjected to the same vectorization, the model evaluates the quality as the inverse of the Euclidean distance between vectors \mathbf{t} , and \mathbf{p} , as

$$D(\mathbf{t}, \mathbf{p}) = \sqrt{\sum_{i=1}^n (t_i - p_i)^2}.$$

$$R(\mathbf{t}, \mathbf{p}) = D(\mathbf{t}, \mathbf{p})^{-1}$$

The reward administered, R , acts as a simple way to calculate the expected similarity between two sequences, which can be re-evaluated upon continuously with a sliding window over the most recent 2,560 datapoints.

2.4 Implementation

Three autoencoder models were implemented in PyTorch based on this task: a 1-Layer LSTM, 3-Layer LSTM, and 2-Headed Transformer. All models took a sequence of size 128 as an input and embedded the sequences in a 16-dimensional space. The performance of each model was compared with the performance of Cosine-Similarity, the most simple and widely used algorithm for comparison.

2.5 Software Design

In addition to optimizing the output of our model, it is essential to consider the continuous datastream output from GCaMP sensors, particularly in terms of performance and safety. For this, a novel software package titled *Rasa* was developed in Rust as a high-performance, multi-threaded, visual solution to integrate analytical tools with memory-safe and accurate data-processing. The previously described Arduino-Photosensor circuit is taken as an input, allowing data to be transferred to the computational software with a USB output bus at very high baud rate of 115,200. The majority of the initial dataset was generated with a provisional Python implementation to log and visualize data that was constructed prior to the beginning of this study. A key goal for the data processing of this study was optimizing this provisional model such that data could be read and analyzed with a much lower latency, allowing stimulation to be delivered closer to the observation of the neural signal. The software was designed with thread-safe software infrastructure to separate the visualization, input, and analytical components with synchronous channels of two-sided deques from which the most

recent data points read from the Arduino can be served to the desired consumer.

3 Results

3.1 Model Evaluation




Figure 3: 1,325 Latent space vectors mapped to PCA embeddings, further mapped to t-SNE embeddings, clustered with K-Means to visualize continuous network of signals and strong peak outliers. Two regions selected from flat and peaked areas of distribution. Blue: right cortex channel. Red: left cortex channel. Vertical offset is arbitrary and normalized as described.

In order to evaluate the quality of each model, the raw embedding vectors were clustered initially with the SciPy implementation of the PCA algorithm to form an initial 2D representation. The PCA points were then passed through a t-SNE algorithm and finally clustered with K-Means. The t-SNE visualization shows a largely continuous and evenly distributed grouping of possible signals with the exception of explicit, correlated spikes, which form a separate cluster. Correlated spikes between the left and right motor cortex are one of the most recognizable common features that occur endogenously and were selected as the point of comparison for the models. This t-SNE cluster was deemed a superset of all peaks observed in the evaluation split of the dataset, as it included many other extreme features that were not strictly peaks. Due to the observed embedding-space distance between points

within the cluster, this inconsistency is believed to be a result of the t-SNE clustering which biases against individual outlying points, as opposed to a flaw in the embeddings. Within this dataset, sequences empirically considered to demonstrate a correlated peak formed a testing dataset by which the models could be evaluated for correctness.

Model	Parameters	Correct (%)	$\frac{D(\text{peak})^{-1}}{D(\text{rest})^{-1}}$	Latency (ms)
Cos. Similarity	0	62.5	11.02	3.72
1-Layer LSTM	12961	92.4	691.1	2.72
3-Layer LSTM	29857	90.7	638.8	4.08
Transformer	277297	98.7	827.2	4.77

We compared the specifications of each tested model for their performance and accuracy. Furthermore, the power of the detected peaks relative to resting stream was considered by

$$\frac{D(\text{peak})^{-1}}{D(\text{rest})^{-1}}$$

which serves to represent the ratio between the inverse distance of a desirable feature, represented as a peak of D^{-1} in the reward graph, and the resting average distance of the signals compared to the reward signal. A larger value for this implies that the peaks are more separated from non-peaks, showing a stronger confidence in the detection.

Notably the latency for each system was approximately equal. Latency was averaged between 5000 samples per model, performed on a Google Colab V100 GPU instance. The latency of all models was considered within tolerance for use for real-time neurofeedback. Despite the large parameter size and comparatively high latency, a transformer model was chosen as the ideal design due to its high accuracy.

3.2 Software Evaluation

To compare the effectiveness of the novel software, both the new and old software were tested after five minutes of idling, eliminating cold-start bias and exemplifying the long-term




Figure 4: Latency after 5 minutes by software. $n = 100, p = 2.6 \cdot 10^{-54}$.

performance benefits of optimized data structures provided by the new implementation. In order to evaluate the latency of the software, an artificial signal would be supplied to the datastream buffer of the respective software every 100 milliseconds by a separate independent timer thread. Following this, the active thread would exhaustively search for the signal among the 128 most recent entries, recording the difference in time between the most recent signal and the beginning of the search. Bias was eliminated by controlling for the search time in Rust ($n = 5; \mu = 0.133$ ms) and Python ($n = 5; \mu = 0.805$ ms).

The distribution of latencies observed from the experiment were evaluated with a two-sample T-Test under the hypothesis that the new method has a lower latency than the previous method. The test evaluated to a p-value significant past $\alpha = 0.0001$, indicating a strong confidence in the improved performance of the novel software.

3.3 *In Vivo* Evaluation

The model and software were integrated into the software package Rasa and run on the McGovern Institute 9.4T fMRI Imaging Setup for an EPI scan under neurofeedback. Feedback was administered as described with the use of a deep-brain electrode as a response to desirable motor cortex activity. For simplicity, we rewarded correlated peaks of electrical ac-




Figure 5: Green: Sum of rewarded values under stimulation based on new criterion. Blue: Sum under no stimulation. $p = 0.0021$

tivity in the two neuronal ensembles considered by the two optical fibers. For the temporarily paralyzed rat, the neural signals and neurofeedback were monitored for two subsequent 20-minute sessions, the first of which generated a "naive" dataset under no stimulation and the second of which was performed with neurofeedback monitoring. During the second session, the model and software achieved a correlated peak with D^{-1} above the arbitrary target threshold 15 times and was stimulated 13 times (discrete peaks separated from one another by less than 16 seconds removed were noted, but not rewarded in order to prevent overstimulation and preserve the health of the animal). Given an arbitrary threshold for inverse distance at value of 300, the sum of every inverse distance greater than that threshold can be thought of as evaluation of both the frequency and intensity of desirable behavior in the data. The value of this threshold is a hyperparameter and its flexibility serves as a principal advantage of this method. The threshold can be constant or changing, depending on the target signal. As such, its value is controllable inside of the Rasa software, allowing researchers to control the frequency and lenience with which reward to be administered.

This training session provides strong evidence that the use of stimulation under the novel system provides some benefit in improving the quantity of rewardable events based on a vector embedding. However, due to the small training size and lack of controls, it did not show strong evidence of improvement over the previous method ($p = 0.20 +$ No unbiased

comparison).

4 Discussion

The method showed promise in our initial experiments as a method of defining neurofeedback qualification. Unfortunately, due to time constraints and the high demand for use of the 9.4T fMRI, this experiment was only able to be performed once. During the two-week *in vivo* portion of this study, there were no opportunities to use the machine during weekdays between the hours of 8AM-5PM, which were the only times available due to laws involving supervision of minors in bio-hazardous laboratories. This prevented the execution of the initial experimental design, which was to test the novel method’s unique ability to encode more complex signals by attempting to reinforce the embedding of a non-correlated peak, a signal that does not generally occur endogenously within rats (constituted 1 of 13,709 sequences in unabridged dataset). This more complex task represents the fundamental advantage of the autoencoder model and loss mechanism. Despite this, the full experiment is scheduled now on the fMRI/Photosensor setup for September-October 2023, likely occurring over the course of 5-8 multi-hour training sessions to account for the increased signal complexity and additional robustness checks. Even though the central experiment was not executed, the Rasa software system and use of ML-based signal analysis for simple signals is beginning to underscore the wider neurofeedback procedure outside the scope of this experiment in isolation.

5 Conclusions

This paper outlines the developmental process of the software system Rasa and its constituent transformer model, which represents a novel technique for reinforcing cognition

modulation based on closed-loop neural signals. We succeeded in demonstrating the performance effectiveness of the software, which has served to aid in more robust data analysis, and improving the ability to detect subtler, and shorter neuronal signals. Furthermore, we developed three unsupervised models that demonstrated the ability to identify arbitrary neural signals and reward them based on the literal positioning of the embedding vectors.

6 Code Availability

The source code for Rasa, as described in this paper, is available on GitHub at <https://github.com/quothonney/rasa>.

7 Acknowledgements

I would like to thank Kevin Chung, Dr. Alan Jasanoff, and the Jasanoff Lab for mentoring me on this project. My tutor Sepehr Asgari and teaching assistant Allen Lin have been unwavering in their support. I am grateful to the Center for Excellence in Education and the Massachusetts Institute of Technology for sponsoring and hosting the Research Science Institute. Finally, I would like to especially express my gratitude to the McGovern Institute for Brain Research for the use of its highly specialized equipment and the Department of Defense for its generous sponsorship.

References

- [1] J. Millan. On the need for on-line learning in brain-computer interfaces. In *2004 IEEE International Joint Conference on Neural Networks (IEEE Cat. No.04CH37541)*, volume 4, pages 2877–2882 vol.4, July 2004.
- [2] S. Enriquez-Geppert, R. J. Huster, and C. S. Herrmann. EEG-Neurofeedback as a tool to modulate cognition and behavior: A review tutorial. *Front Hum Neurosci*, 11:51, Feb. 2017.
- [3] H. Marzbani, H. R. Marateb, and M. Mansourian. Neurofeedback: A comprehensive review on system design, methodology and clinical applications. *Basic Clin Neurosci*, 7(2):143–158, Apr. 2016.
- [4] T.-W. Chen, T. J. Wardill, Y. Sun, S. R. Pulver, S. L. Renninger, A. Baohan, E. R. Schreiter, R. A. Kerr, M. B. Orger, V. Jayaraman, L. L. Looger, K. Svoboda, and D. S. Kim. Ultrasensitive fluorescent proteins for imaging neuronal activity. *Nature*, 499(7458):295–300, July 2013.
- [5] G. Karvat, A. Schneider, M. Alyahyay, F. Steenbergen, M. Tangermann, and I. Diester. Real-time detection of neural oscillation bursts allows behaviourally relevant neurofeedback. *Communications Biology*, 3(1), Feb 2020.
- [6] R. J. Schafer and T. Moore. Selective attention from voluntary control of neurons in prefrontal cortex. *Science*, 332(6037):1568–1571, 2011.
- [7] K. B. Clancy, A. C. Koralek, R. M. Costa, D. E. Feldman, and J. M. Carmena. Volitional modulation of optically recorded calcium signals during neuroprosthetic learning. *Nature neuroscience*, 17(6):807–809, 2014.
- [8] Z. Liang, Y. Ma, G. D. Watson, and N. Zhang. Simultaneous gcamp6-based fiber photometry and fmri in rats. *Journal of Neuroscience Methods*, 289:31–38, 2017.
- [9] S. H. Scott. Optimal feedback control and the neural basis of volitional motor control. *Nature Reviews Neuroscience*, 5(7):532–545, 2004.
- [10] R. Nieuwenhuys, L. M. Geeraedts, and J. G. Veening. The medial forebrain bundle of the rat. i. general introduction. *Journal of Comparative Neurology*, 206(1):49–81, 1982.
- [11] D. P. Kingma, M. Welling, et al. An introduction to variational autoencoders. *Foundations and Trends® in Machine Learning*, 12(4):307–392, 2019.